

PRACTICAL PROJECTS WITH mongoDB



WHO AM I?

Alex Sharp

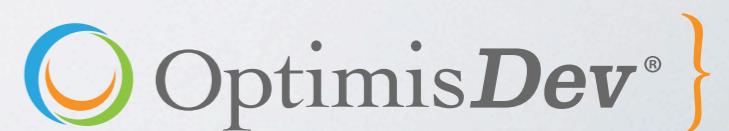
Lead Developer at OptimisDev



@ajsharp (on the twitters)

alexjsharp.com (on the 'tubes)

github.com/ajsharp (on the 'hubs)



MAJOR POINTS

- Brief intro to MongoDB
- Practical Projects
- Making the case for Mongo (time permitting)

BRIEF INTRO

WHAT IS MONGODB?

mongodb is a high-performance, schema-less,
document-oriented database



STRENGTHS

SCHEMA-LESS

SCHEMA-LESS

- Great for rapid, iterative, agile development

SCHEMA-LESS

- Great for rapid, iterative, agile development
- Makes things possible that in an rdbms are either
 - a. impossibly hard
 - b. virtually impossible
 - c. way harder than they should be

DOCUMENT-ORIENTED

DOCUMENT-ORIENTED

- Mongo stores *documents*, not rows
 - documents are stored as binary json

DOCUMENT-ORIENTED

- Mongo stores *documents*, not rows
 - documents are stored as binary json
- Allows for a rich query syntax

DOCUMENT-ORIENTED

- Mongo stores *documents*, not rows
 - documents are stored as binary json
- Allows for a rich query syntax
- Embedded documents eliminate many modeling headaches

BIG DATA

BIG DATA

- Built to scale horizontally into the cloudz

BIG DATA

- Built to scale horizontally into the cloudz
- Auto-sharding
 - set a shard-key, a cluster of nodes, go
 - still in alpha, but production-ready soon

FAST WRITES

FAST WRITES

- Mongo writes are “fire-and-forget”
 - you can get a response with `getLastError()`

FAST WRITES

- Mongo writes are “fire-and-forget”
 - you can get a response with `getLastError()`
- In-place updates

FAST WRITES

- Mongo writes are “fire-and-forget”
 - you can get a response with `getLastError()`
- In-place updates
- Lot’s of db commands for fast updates

AGGREGATION

Map/reduce for aggregation

READS FAST TOO

Optimize reads with indexes, just like an rdbms

WEAKNESSES

JOINS

nope.

MULTI-DOCUMENT TRANSACTIONS

no-sir-ee.

RELATIONAL INTEGRITY

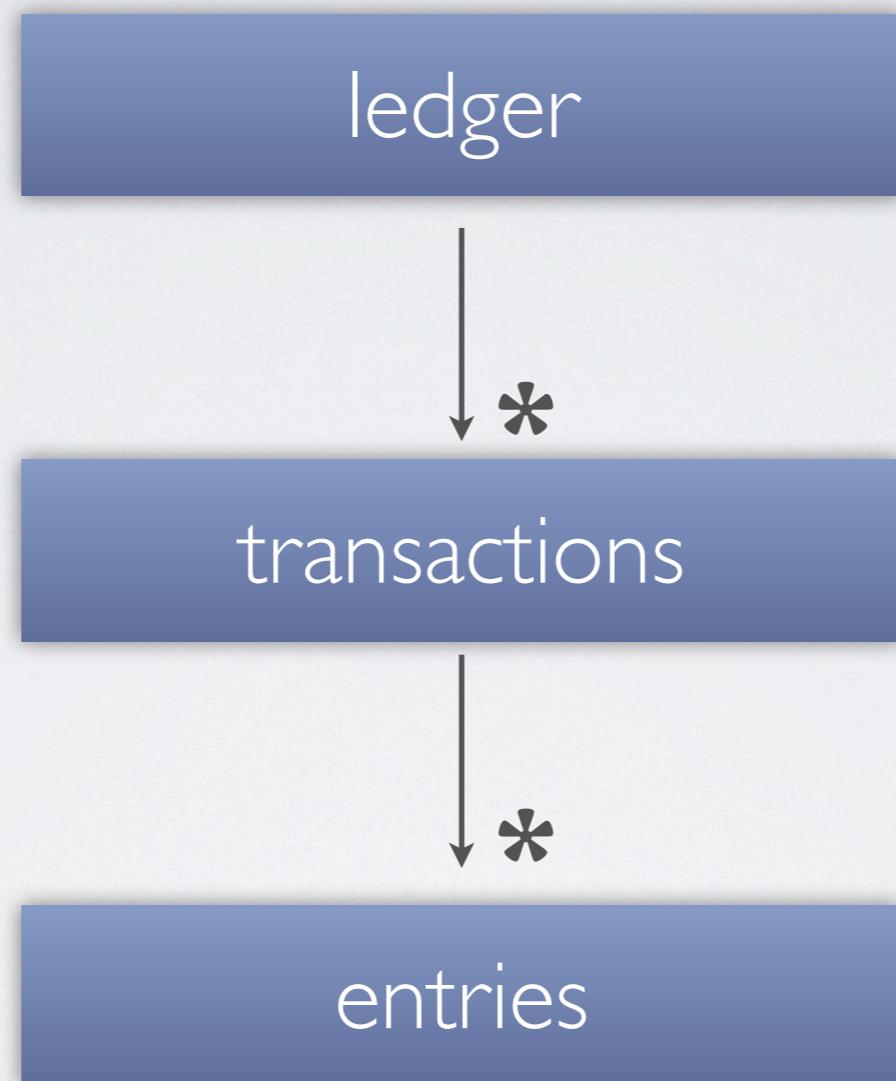
not a chance.

PRACTICAL PROJECTS

1. Accounting Application
2. Capped Collection Logging
3. Blogging app
4. Reporting app

GENERAL LEDGER ACCOUNTING APPLICATION

THE OBJECT MODEL



THE OBJECT MODEL

#	Credits	Debits
1	{ :account => "Cash", :amount => 100.00 }	{ :account => "Notes Pay.", :amount => 100.00 }
2	{ :account => "A/R", :amount => 25.00 }	{ :account => "Gross Revenue", :amount => 25.00 }

Line item { Ledger Entries }

Line item { }

THE OBJECT MODEL

THE OBJECT MODEL

- Each ledger *line item* belongs to a *ledger*

THE OBJECT MODEL

- Each ledger *line item* belongs to a *ledger*
- Each ledger *line item* has two ledger entries
 - must have at least one credit and one debit
 - credits and debits must balance

THE OBJECT MODEL

- Each ledger *line item* belongs to a *ledger*
- Each ledger *line item* has two ledger entries
 - must have at least one credit and one debit
 - credits and debits must balance
- These objects must be transactional

SQL-BASED OBJECT MODEL

```
@debit = Entry.new(:account => 'Cash',  
:amount => 100.00, :type => 'credit')
```

```
@credit = Entry.new(:account => 'Notes Pay.',  
:amount => 100.00, :type => 'debit')
```

```
@line_item = LineItem.new :ledger_id => 1, :entries => [@debit, @credit]
```

SQL-BASED OBJECT MODEL

```
@debit = Entry.new(:account => 'Cash',
:amount => 100.00, :type => 'credit')

@credit = Entry.new(:account => 'Notes Pay.', 
:amount => 100.00, :type => 'debit')

@line_item = LineItem.new :ledger_id => 1, :entries => [@debit, @credit]
```

In a SQL schema, we need a database transaction to ensure multi-row atomicity

SQL-BASED OBJECT MODEL

```
@debit = Entry.new(:account => 'Cash',
:amount => 100.00, :type => 'credit')

@credit = Entry.new(:account => 'Notes Pay.', 
:amount => 100.00, :type => 'debit')

@line_item = LineItem.new :ledger_id => 1, :entries => [@debit, @credit]
```

Remember: we have to create 3 new records in
the database here

MONGO-FIED MODELING

```
@line_item = LineItem.new(:ledger_id => 1,  
:entries => [  
  {:account => 'Cash', :type => 'debit', :amount => 100},  
  {:account => 'Notes pay.', :type => 'debit', :amount => 100}  
]  
)
```

MONGO-FIED MODELING

```
@line_item = LineItem.new(:ledger_id => 1,  
:entries => [  
  {:account => 'Cash', :type => 'debit', :amount => 100},  
  {:account => 'Notes pay.', :type => 'debit', :amount => 100}  
]  
)
```

This is the perfect case for embedded documents.

MONGO-FIED MODELING

```
@line_item = LineItem.new(:ledger_id => 1,  
:entries => [  
  {:account => 'Cash', :type => 'debit', :amount => 100},  
  {:account => 'Notes pay.', :type => 'debit', :amount => 100}  
]  
)
```

We will never have more than a few entries (usually two)

MONGO-FIED MODELING

```
@line_item = LineItem.new(:ledger_id => 1,  
:entries => [  
  {:account => 'Cash', :type => 'debit', :amount => 100},  
  {:account => 'Notes pay.', :type => 'debit', :amount => 100}  
]  
)
```

Embedded docs are perfect for “contains” many relationships

MONGO-FIED MODELING

```
@line_item = LineItem.new(:ledger_id => 1,  
:entries => [  
  {:account => 'Cash', :type => 'debit', :amount => 100},  
  {:account => 'Notes pay.', :type => 'debit', :amount => 100}  
]  
)
```

DB-level transactions no-longer needed.
Mongo has single document atomicity.

LESSONS

LESSONS

- Simplified, de-normalized data model
 - Objects modeled differently in Mongo than SQL

LESSONS

- Simplified, de-normalized data model
 - Objects modeled differently in Mongo than SQL
- Simpler data model + no transactions = WIN

LOGGING WITH CAPPED COLLECTIONS

BASELESS STATISTIC

95% of the time logs are NEVER used.

MOST LOGS ARE ONLY USED
WHEN SOMETHING GOES
WRONG

CAPPED COLLECTIONS

- Fixed-sized, limited operation, auto age-out collections (kinda like memcached, except persistent)
- Fixed insertion order
- Super fast (faster than normal writes)
- Ideal for logging and caching

GREAT. WHAT'S THAT MEAN?

We can log additional pieces of arbitrary data effortlessly due to Mongo's schemaless nature.

THIS IS AWESOME

QUERY. ANALYZE. PROFIT.

ALSO A REALLY HANDY
TROUBLESHOOTING TOOL

User-reported errors are difficult to
diagnose

Wouldn't it be useful to see the complete click-path?

BUNYAN

<http://github.com/ajsharp/bunyan>

Thin ruby layer around a
MongoDB capped collection

BUNYAN

<http://github.com/ajsharp/bunyan>

Still needs a middleware api.

Want to contribute?

```
require 'bunyan'
```

```
Bunyan::Logger.configure do |c|
  c.database    'my_bunyan_db'
  c.collection "#{{Rails.env}}_bunyan_log"
  c.size        1073741824 # 1.gigabyte
end
```

PAPER MILL

<http://github.com/ajsharp/papermill>

SINATRA FRONT-END TO BUNYAN

BLOGGING APPLICATION

BLOGGING APPLICATION

Much easier to model with Mongo than a relational database



BLOGGING APPLICATION

A post has an author



BLOGGING APPLICATION

A post has an author

A post has many tags



BLOGGING APPLICATION

A post has an author

A post has many tags

A post has many comments



BLOGGING APPLICATION

A post has an author

A post has many tags

A post has many comments

Instead of JOINing separate tables,
we can use embedded documents.



MONGOMAPPER



MONGOMAPPER

- MongoDB “ORM” developed by John Nunemaker



MONGOMAPPER

- MongoDB “ORM” developed by John Nunemaker
- Very similar syntax to DataMapper



MONGOMAPPER

- MongoDB “ORM” developed by John Nunemaker
- Very similar syntax to DataMapper
 - Declarative rather than inheritance-based



MONGOMAPPER

- MongoDB “ORM” developed by John Nunemaker
- Very similar syntax to DataMapper
 - Declarative rather than inheritance-based
- Very easy to drop into rails



MONGOMAPPER

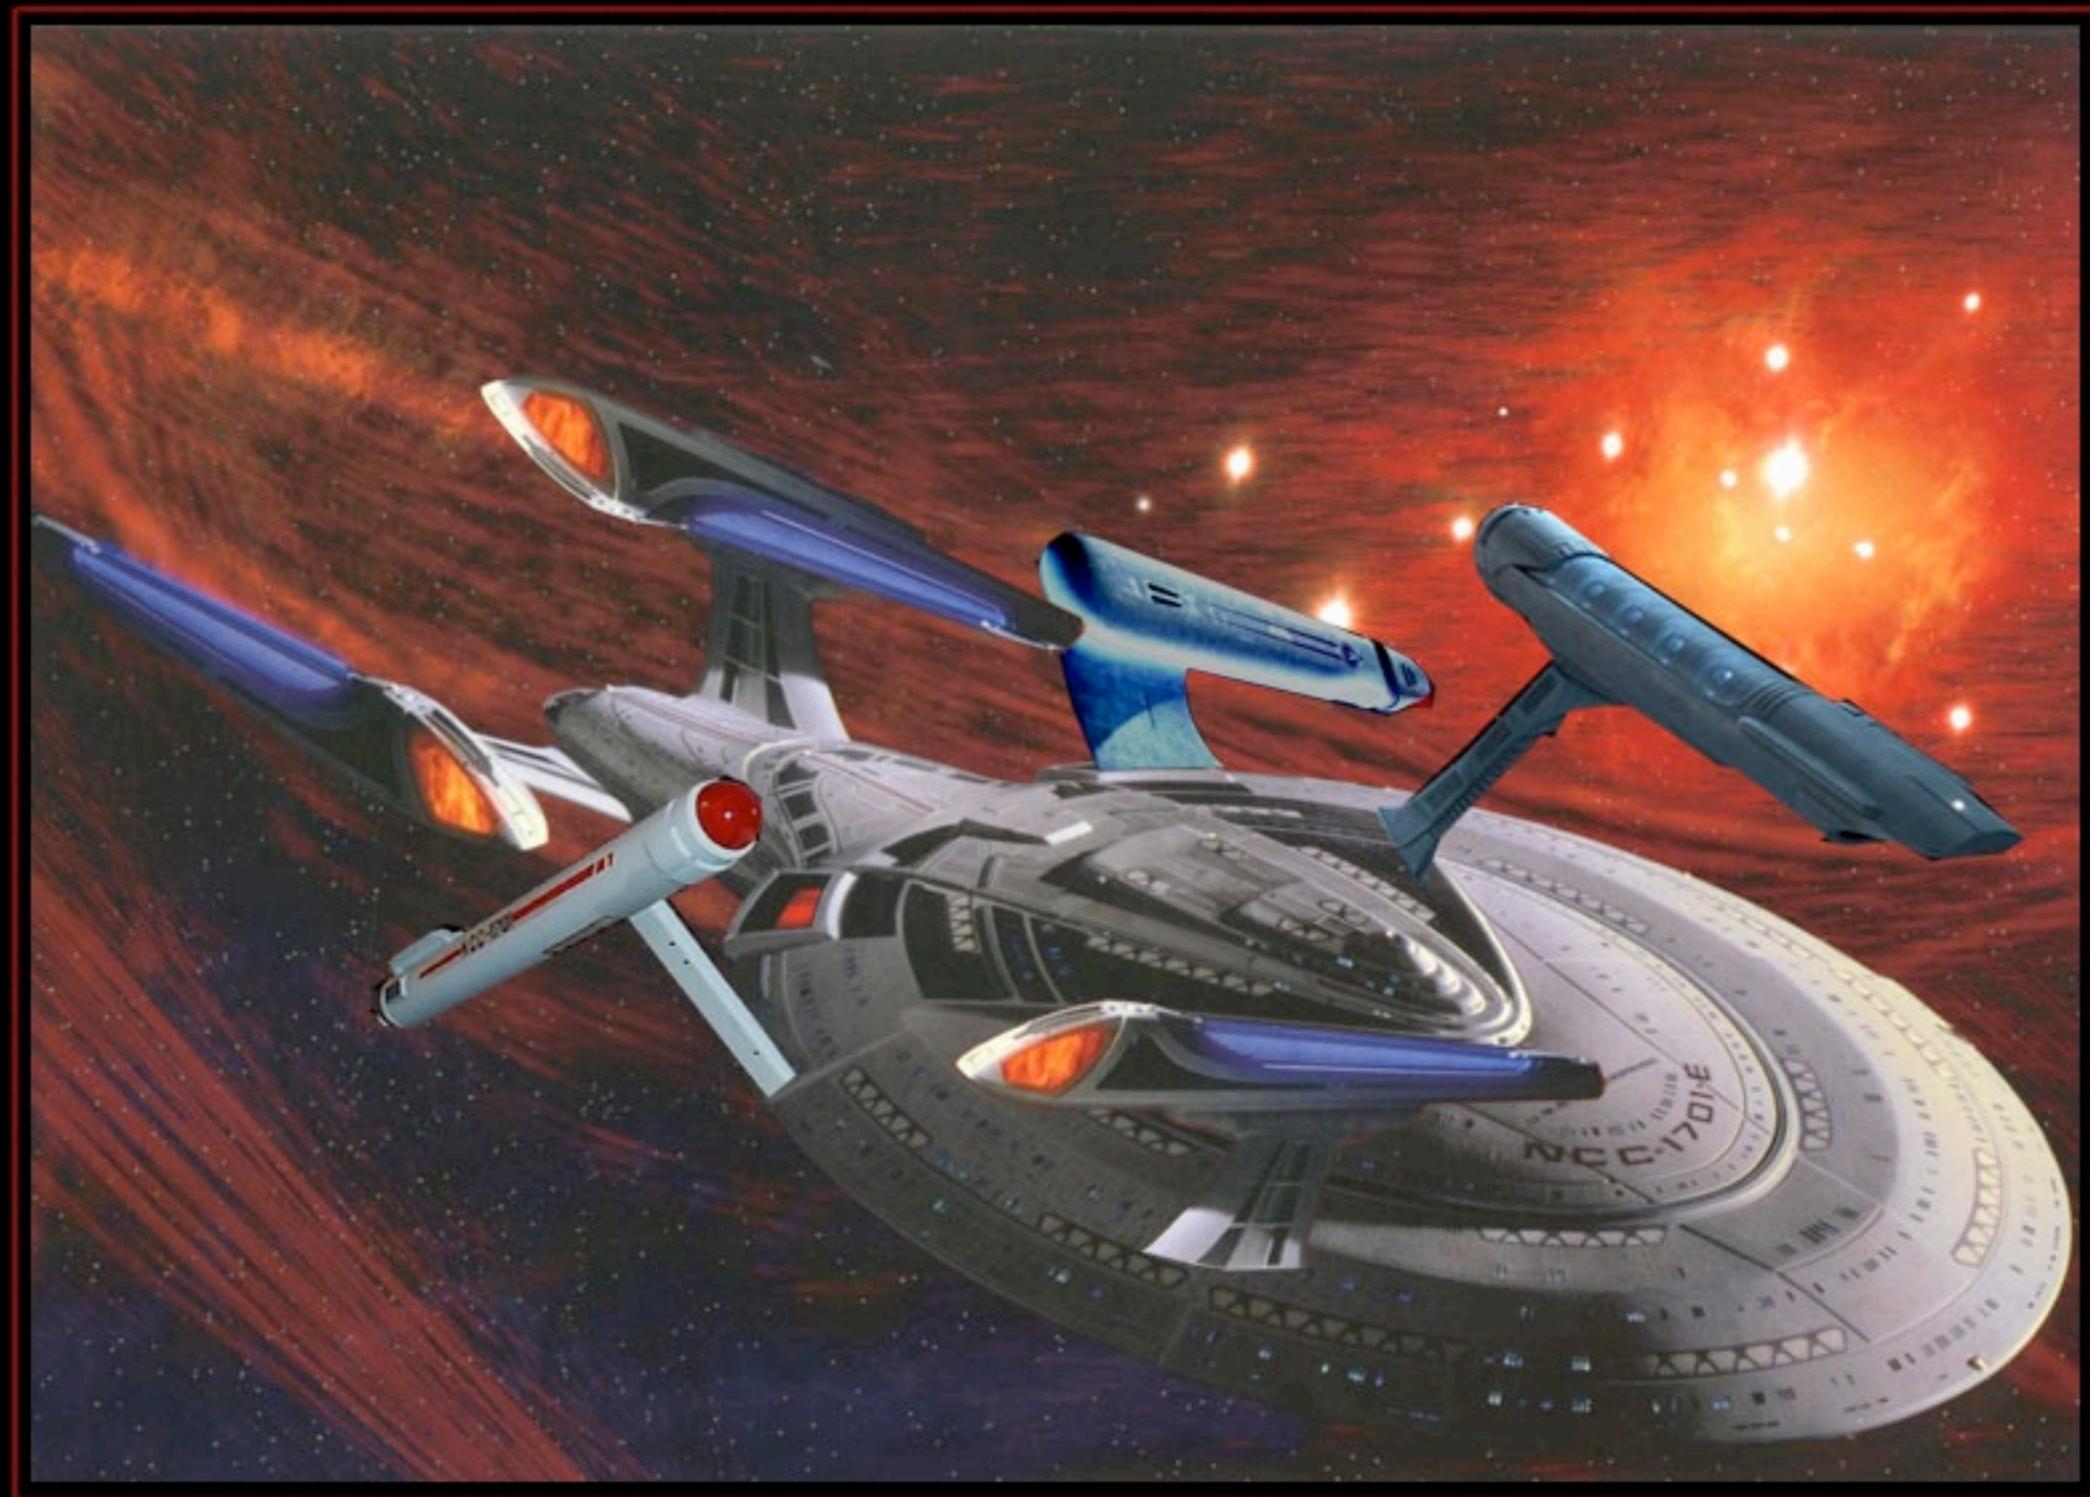
```
require 'mongo'
connection = Mongo::Connection.new.db("#{Rails.env}_app")

class Post
  include MongoMapper::Document

  belongs_to :author, :class_name => "User"

  key :title,           String,  :required => true
  key :body,            String,  :required => true
  key :author_id,       Integer, :required => true
  key :published_at,    Time
  key :published,       Boolean, :default => false
  key :tags,             Array,   :default => []
  timestamps!
end
```

REPORTING APP



ENTERPRISE SOFTWARE

It's really complicated, it took forever to build, and it cost a whole lot.

*It **must** be better.*

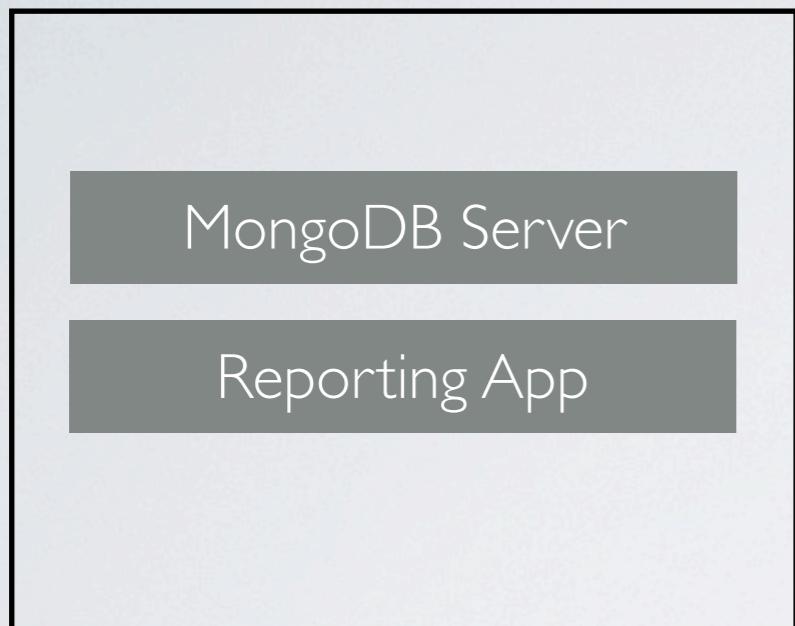
REPORTING IS HARD

- Synchronization
 - Data, application API, schema
- Schema synchronization is the hard one

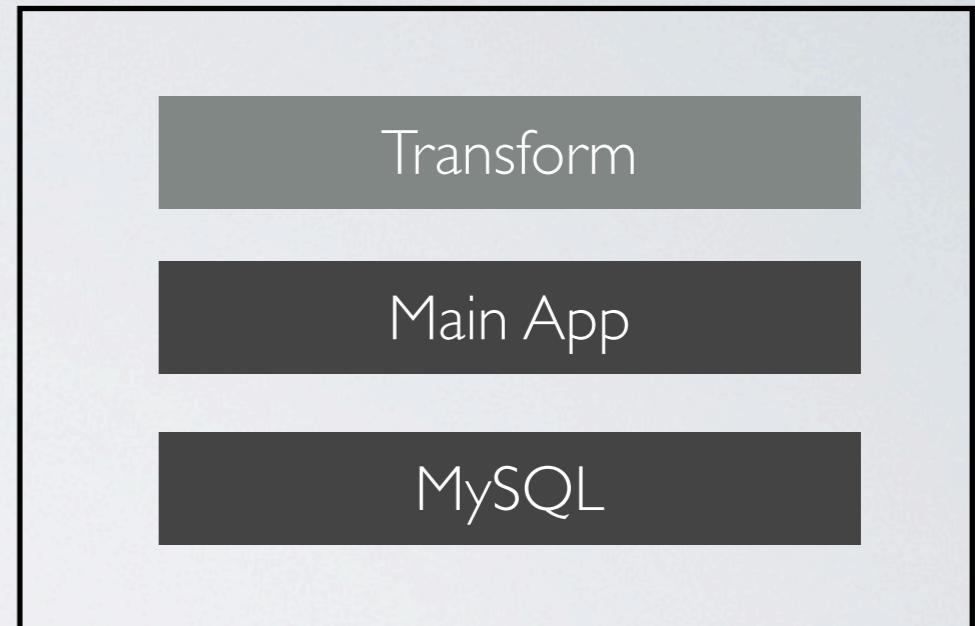
WHY MONGO?

- Aggregation support with Mongo's map/reduce support
- Mongo is good (and getting better) at handling big data sets
- Schemaless is perfect for a flattened data set

Amazon EC2



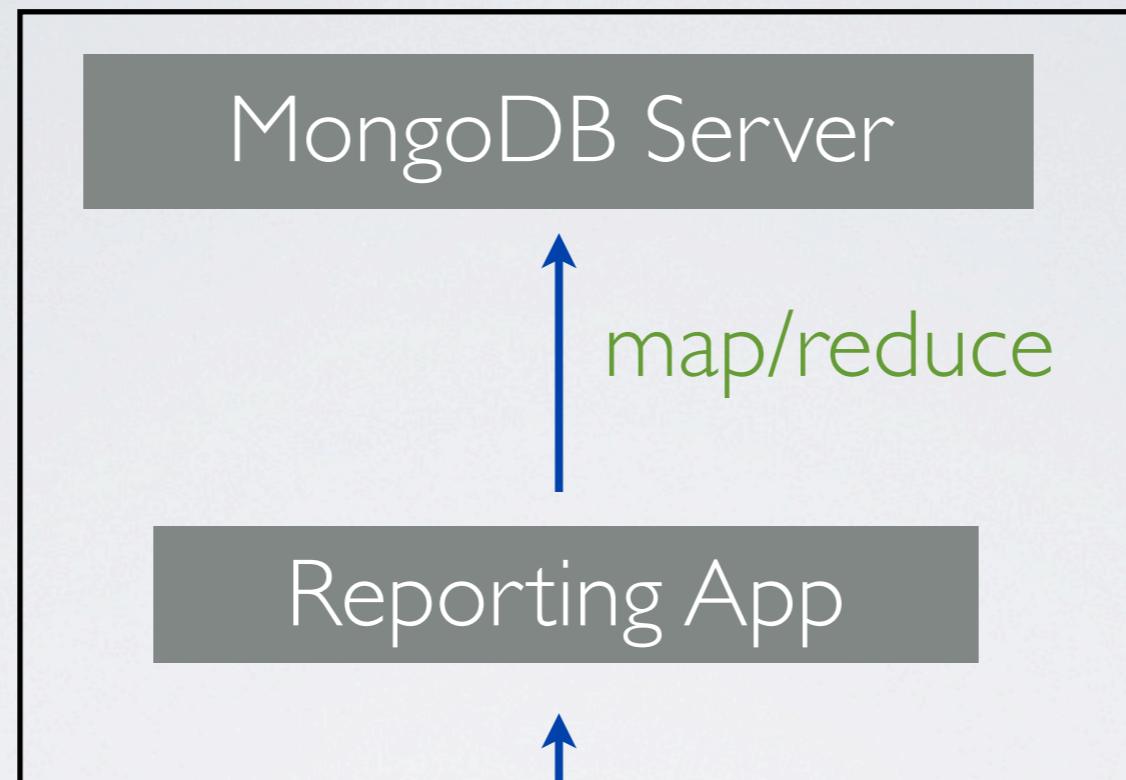
Rackspace



WAN

CLIENT REQUESTS

Amazon EC2



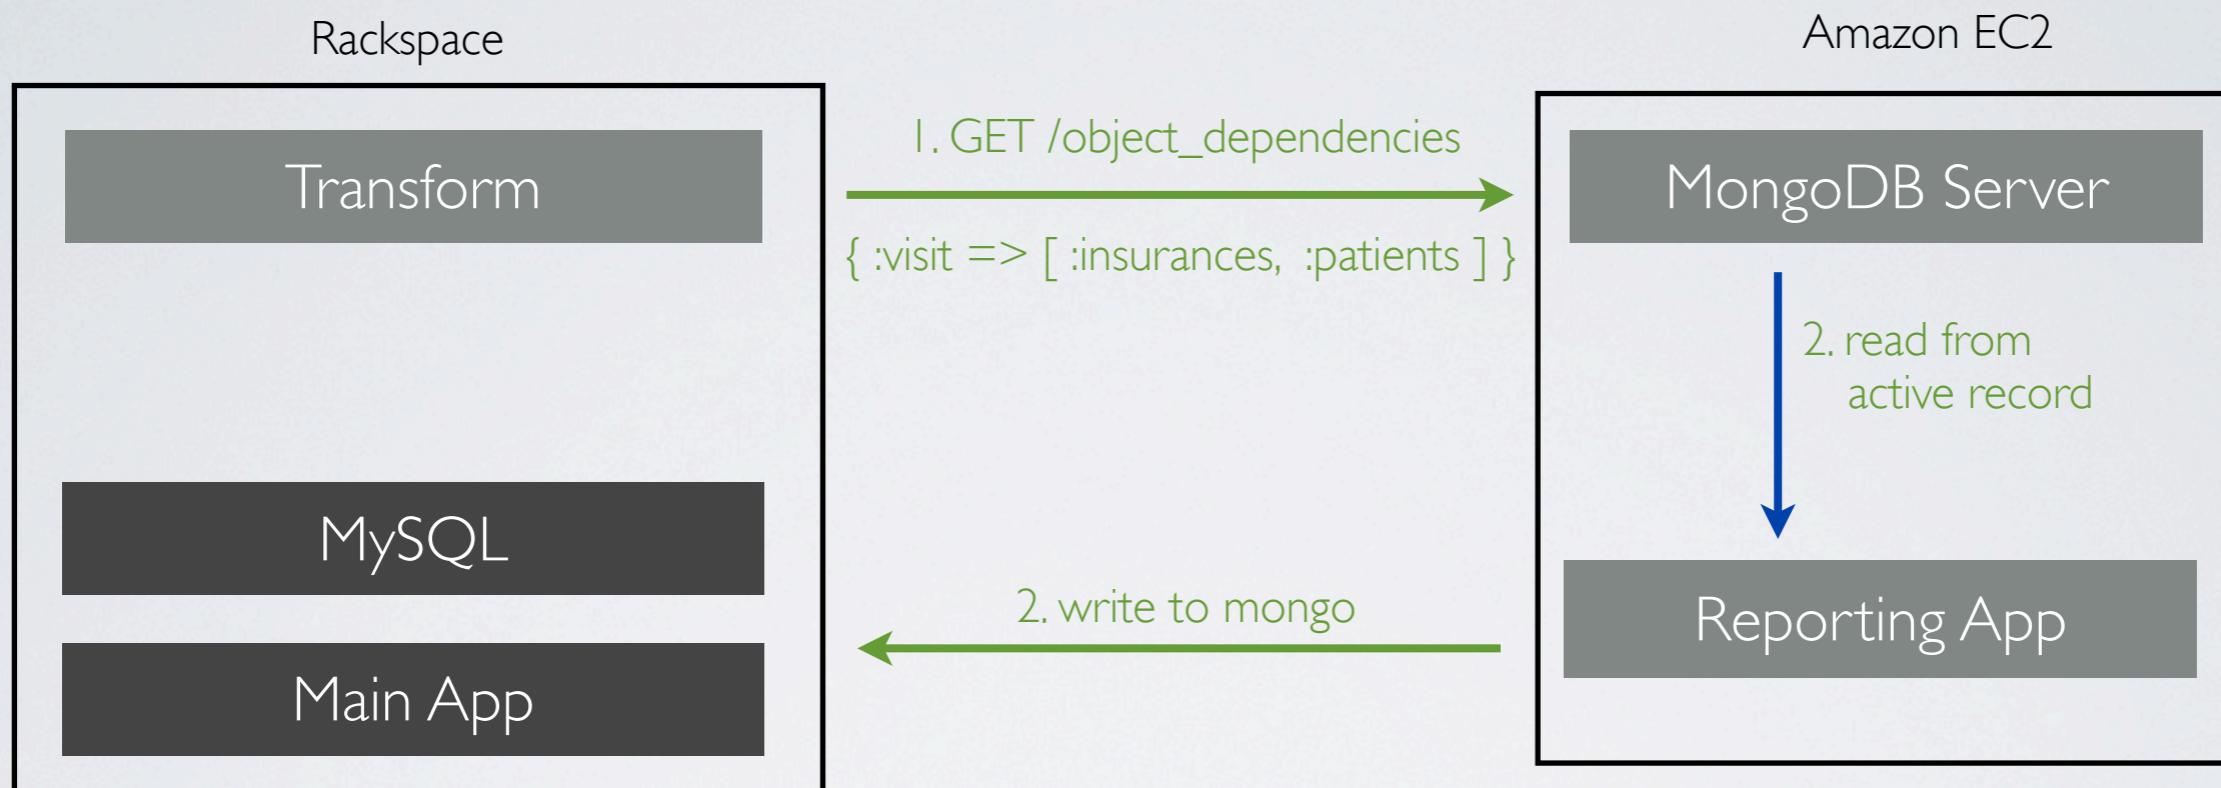
Client request `GET /reports/some_report`

User requests

TRANFORMATION

Data needs to extracted from the main app,
transformed (i.e. flattened) into some other form,
and loaded into the reporting app for
aggregation

TRANSFORMATION



MAKING THE CASE

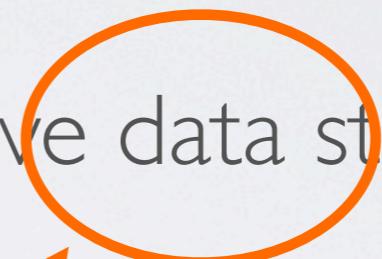
THE PROBLEMS OF SQL

A BRIEF HISTORY OF SQL

- Developed at IBM in early 1970's.
- Designed to manipulate and retrieve data stored in relational databases

A BRIEF HISTORY OF SQL

- Developed at IBM in early 1970's.
- Designed to manipulate and retrieve data stored in relational databases



this is a problem

WHY??

WE DON'T WORK WITH
DATA...

WE WORK WITH OBJECTS

WE DON'T CARE
ABOUT
STORING DATA

WE CARE ABOUT
PERSISTING STATE

DATA != OBJECTS

THEREFORE...

RELATIONAL DBS ARE AN
ANTIQUATED TOOL FOR OUR
NEEDS

OK, SO WHAT?

SQL schemas are designed for storing and querying *data*, not *persisting objects*.

To reconcile this
mismatch, we have
ORM's

Object Relational *Mappers*

OK, SO WHAT?

We need ORMs to bridge the gap
(map) between our native Ruby
object model and a SQL schema

We're forced to create *relationships* for *data*
when what we really want is *properties* for our
objects.

THIS IS NOT EFFICIENT

```
@alex = Person.new(  
  :name => "alex",  
  :stalkings => [  
    Friend.new("Jim"),  
    Friend.new("Bob")]  
)
```

Native ruby object

```
<Person:0x10017d030 @name="alex",
@stalkings=
  [#<Friend:0x10017d0a8 @name="Jim">,
  #<Friend:0x10017d058 @name="Bob">
]>
```

JSON/Mongo Representation

```
@alex.to_json
{ name: "alex",
  stalkings: [{ name: "Jim" }, { name: "Bob" }]
}
```

SQL Schema Representation

people:

- name

stalkings:

- name
- stalker_id

SQL Schema Representation

people:

- name

stalkings:

- name
- stalker_id

What!?



SQL Schema Representation

people:

- name

stalkings:

- name

- stalker_id

what's a stalker_id?



SQL Schema Representation

people:

- name

stalkings:

- name
- stalker_id

this is our sql mapping



SQL Schema Representation

people:

- name

stalkings:

- name
- stalker_id

this is the
“pointless” join



Ruby -> JSON -> SQL

Ruby



JSON



SQL

```
<Person:0x10017d030 @name="alex",
@stalkings=
[#<Friend:0x10017d0a8 @name="Jim">,
 #<Friend:0x10017d058 @name="Bob">
]>

@alex.to_json
{ name: "alex",
  stalkings: [{ name: "Jim" }, { name: "Bob" }]
}

people:
  - name

stalkings:
  - name
  - stalker_id
```

Ruby -> JSON -> SQL

Ruby

```
<Person:0x10017d030 @name="alex",
@stalkings=
[#<Friend:0x10017d0a8 @name="Jim">,
 #<Friend:0x10017d058 @name="Bob">
]>

@alex.to_json
{ name: "alex",
  stalkings: [{ name: "Jim" }, { name: "Bob" }]
}
```

JSON

SQL

```
people:
- name

stalkings:
- name
- stalker_id
```

Feels like we're having
to work too hard here

You're probably thinking...

“Listen GUY, SQL isn’t *that* bad”

Maybe.

But we can do much, much better.

NEEDS FOR A PERSISTENCE LAYER:

NEEDS FOR A PERSISTENCE LAYER:

- I. To persist the native state of our objects

NEEDS FOR A PERSISTENCE LAYER:

1. To persist the native state of our objects
2. Should NOT interfere w/ application development

NEEDS FOR A PERSISTENCE LAYER:

1. To persist the native state of our objects
2. Should NOT interfere w/ application development
3. Provides features necessary to build modern web apps

NEXT STEPS

USING MONGO W/ RUBY

- Ruby mongo driver
- MongoMapper (github.com/jnunemaker/mongomapper)
- Mongoid (github.com/durran/mongoid)
- Many other ORM/ODM's
 - moneta (github.com/wycats/moneta)

MONGOMAPPER

- DataMapper-like API
- good for moving from a sql schema to mongo
- easy to drop into rails
- works with rails 3

MONGOID

- DataMapper-like API
- uses ActiveRecord, so familiar validation syntax
- more opinionated embedded docs
- easy to drop into rails
- rails 2 - mongoid 1.x
- rails 3 - mongoid 2.x

QUESTIONS?

THANKS!